

Security of Wireless Networks

{thgoebel}@ethz.ch

ETH Zürich, HS 2020

This is a summary for the course *Security of Wireless Networks (SOWN)* at ETH Zurich.

This summary is created during the autumn semester 2020. But due to the few changes in syllabus content in the past we have reason to believe that it is also relevant beyond that very semester.

We do not guarantee correctness or completeness, nor is this document endorsed by the lecturers. Feel free to point out any erratas.

Contents

1. Wireless Basics	3
2. Jamming Basics	8
2.1. Jamming Resistant Communication	8
3. Jamming-Resistant Broadcast	11
4. Security of Global Navigation Satellite Systems GNSS	15
4.1. Spoofing Detection and Mitigation	15
5. Secure Distance Measurement	17
6. Security in Critical Transport Infrastructures	21
6.1. Privacy Issues in Aviation	21
6.2. Security Issues in Aviation	22
6.3. Short-term Security Countermeasures	23
6.4. Satellite and Maritime Infrastructures	23
7. Physical Layer-based Security	24
7.1. Channel-based Key Establishment	24
7.2. Friendly Jamming for Confidentiality and Access Control	25
7.3. Broadcast Authentication: Presence Awareness	27
8. Broadcast Authentication and Device Pairing	28
8.1. Broadcast Authentication: Delayed Key Disclosure	28
8.2. Device Pairing	28
9. WiFi Security	31
9.1. WiFi Basics	31
9.2. Basic Manipulations	33
9.3. WiFi Security Standards	33
A. Imprint	36

1. Wireless Basics

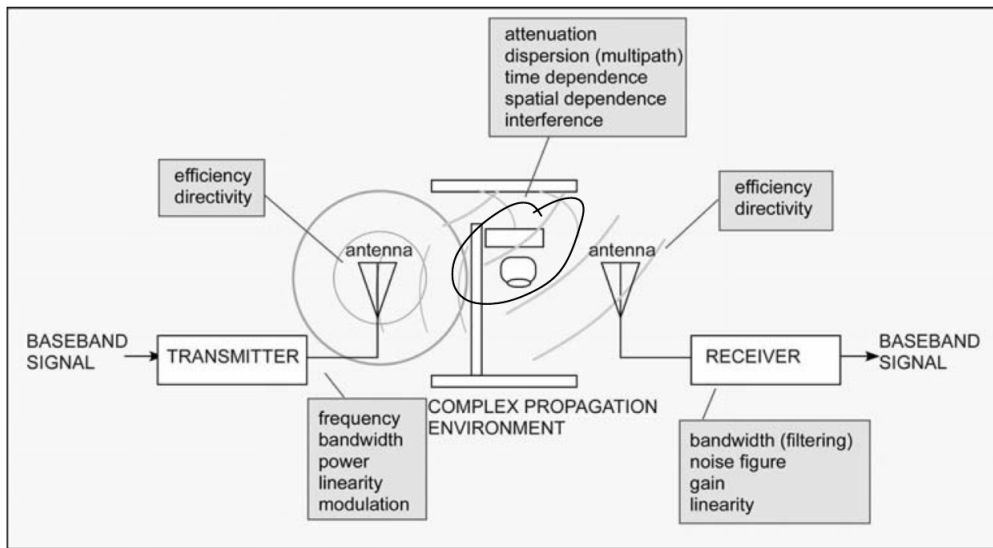


Figure 1: A wireless system, its basic components and characteristic measures

Radio Frequency Signal Electromagnetic radiation, with waves being created in the antenna by an alternating current at the desired frequency. Mathematically described as a function of the time t :

$$v(t) = A \sin(2\pi ft + \phi)$$

with amplitude A , frequency f and phase ϕ . Also recall that the period is $T = \frac{1}{f}$ and the wavelength (distance travelled during one period) is $\lambda = \frac{v}{f}$ (usually $v = c$ speed of light).

Bandwidth The capacity of a communications link to transmit the maximum amount of data from one point to another over a connection in a given amount of time (in bits per second bps). An analogy: The amount of water that can flow through a water pipe.

In other words, the measure of frequency content of the signal. E.g. the human voice contains frequencies in the range from 30 Hz to 10 kHz, and the bandwidth of a single 802.11 channel is 22 MHz.

Note that often the bandwidth of the baseband and that of the carrier (and thus that of the modulated signal) differ! E.g. see spread spectrum techniques (subsection 2.1).

Baseband An original transmission signal that has not been modulated or has been demodulated to its original frequency. I.e. the actual **information signal**. Most telecommunication protocols require baseband signals to be converted, or modulated, to a higher frequency in order to be transmitted over long distances.

Carrier A transmitted electromagnetic pulse or wave at a steady base frequency of alternation on which information can be imposed. Typically a pure sinusoid of a particular frequency and phase that carries the information. Usually the frequency of the carrier is much higher than that of the baseband.

Modulated Signal A carrier that has been loaded or modulated with the information signal.




Figure 2: 2.4 GHz WiFi Channels [Source]

Modulation Process of imposing the baseband onto the carrier. The baseband is used to alter one aspect of the carrier, such as: signal strength (*amplitude modulation AM*), frequency (*frequency modulation FM*), phase (*phase modulation PM*). In other words, one of the values A, f, ϕ in the above equation of the signal is manipulated.

Phase-shift keying PSK Modulation technique varying the phase of the carrier. Used e.g. in WiFi, RFID, Bluetooth. Specific versions include Binary PSK, Quadrature PSK and Differential PSK. Simple example: if the baseband bit is 0 do nothing to the carrier, if it is 1 shift the carrier phase by π .

On-Off-Keying OOK Simple form of amplitude-shift keying ASK. Represents data as the presence (1) or absence (0) of a signal. E.g. Morse code.

I-Q Signal Representation A pair of periodic signals are said to be in ‘quadrature’ when they differ in phase by 90 degrees (e.g. the sine and cosine wave). The ‘in-phase’ or reference signal is referred to as ‘I’ (conventionally cosine), and the signal that is shifted by 90 degrees (in quadrature) is called ‘Q’ (conventionally sine). Used to represent modulations.

Antenna Interface between radio waves in the air and electric alternating currents in a conductor. Types include: omni/dipole, yagi, horn, cantenna.

The directionality of an antenna describes how well it transmits/receives into a particular direction.

- **isotropic** – Theoretical, radiates with the same intensity equally in all directions. Often used as a reference antenna when calculating the gain.
- **omni-directional** – Radiates equally well in all directions in a flat horizontal plane. Most common types in consumer devices.
- **directional** – Radiates best in a given direction by focussing its power. Can thus work with weaker signals than an omni-directional antenna of the same power.

Phased Array Array of fixed antennas where the phase of each signal is dynamically adjusted so that the signal will be in phase for a given direction. Allows *beam steering* towards a specific direction. Possible applications? Can it be used to achieve security (e.g. confidentiality)?

Transmitter/Receiver Converts from digital to analogue, applies modulation and connects to the antenna (and vice versa). Properties: transmitted power, carrier frequency, information bandwidth, modulation type, receiver sensitivity.




Figure 3: Antenna directionality




Figure 4: Beam steering




Figure 5: Signal strength across the channel (ERP = Effective Radiated Power)

Software Defined Radio SDR Flexible, low-cost transmitter/receiver. Implements components (mixer, amplifier, de-/modulator) in software rather than processing the signal in hardware.

Channel equation See Figure 5.

signal strength at the receiver = transm. power + transm. antenna gain – link loss + receiv. antenna gain

Note that in free space the power density of an EM wave obeys the inverse-square law:

$$p \propto \frac{1}{d^2}$$

Receiver sensitivity The weakest signal from which the receiver can still obtain the desired information signal. Depends not just on the antenna gain, but also on other factors such as the noise.

Decibel

- dBm – signal strength in dB / 1 milliwatt mW
- dBW – signal strength in dB / 1 watt W
- dBi – antenna gain in dB / antenna gain of isotropic antenna in dB

Calculating a value in dB:

$$dB(n) = 10 \log_{10}(n) \quad \text{and} \quad dBm(n) = 10 \log_{10}(n/1mW)$$

Power Spectral Density diagram Depicts the power density (in dB) for a range of frequencies. In simple terms, it shows how strong the signal is at a given frequency.

Security Goals Reasons: *security* (integrity, confidentiality, authentication), *regulatory* (personal liability for misuse of one's network access), *safety* (RF-enabled implants).

Just reducing transmission power, hoping that the attacker will be too far away to listen on / send / modify messages, is NOT a solution. In fact, WiFi signals can be received 10 km away, and similarly Bluetooth at 1 km distance (with good, directed equipment).

Example: *passive keyless entry and start systems (PKES)*, i.e. wireless car keys. Wrongly assume communication implies physical proximity (relay attack). Needs: Authenticated proximity verification, message authentication.

2. Jamming Basics

Jamming Entirely preventing or reducing the ability of communicating parties to pass information, either intentionally or unintentionally.

The jamming signal needs to have the same frequency as the modulated signal. If the latter is unknown to the attacker, they thus need to jam a wide bandwidth of frequencies to be successful.

Effectively, jamming is always a power play.

Symbol Carries one or more bit of information, depending on the modulation scheme.

Symbol Jamming Corrupts symbols such that the receiver can EITHER not interpret them OR interprets them incorrectly.

Targeted, low-power jamming of specific symbols is hard!

Communication Jamming Corrupts enough bits that the information cannot be reconstructed any more, despite error correction.

Jamming-to-Signal Ratio J/S = $J - S$, i.e. the difference between the jamming signal and the modulated signal in dB. A ratio ≥ 0 usually results in successful jamming.

Burn-through range Range in which communication still succeeds, despite jamming.

Attacker model

Types: responsive, sweep, random

Actions: jam, insert, modify (= overshadow)

Power to jam/insert/modify: P_j, P_t, P_o

channels to jam/insert/modify: c_j, c_t, c_o

Total strength/power P_T

$$c_j P_j + c_t P_t + c_o P_o \leq P_T$$

2.1. Jamming Resistant Communication

Basic principle If you cannot fight (i.e. have too little power), RUN, HIDE or WAIT. And get an advantage over the attacker: use a shared secret.




Figure 6: DSSS – hiding under the noise

Frequency Hopping Spread Spectrum FHSS Regularly change transmission frequency. The pseudo-random frequency sequence is derived from a shared secret. Sender and receiver **must** be synchronised.

Note that frequency hoppers can be detected and located, simply by looking over time from which direction someone is sending on changing frequencies.

Possible attacks:

- **Partial band jammer:** Distribute jamming power over a subset of all hopping frequencies to achieve $J/S = 0$ at least on that range.
- **Follower jammer:** Detects on which frequency communication occurs and then jams it. Can be protected against by using error codes (since only the final bits will be corrupted).

Direct Sequence Spread Spectrum DSSS Spreads the baseband over a larger bandwidth using a shared secret (narrowband to broadband).

Since the transmission power remains the same, the power density at any given frequency decreases. Thus the spread signal can effectively “hide under the noise” (Figure 6).

To spread over more frequencies, we need a higher symbol/bit rate. To achieve this the information signal is multiplied with a high-frequency pseudorandom sequence called **chips** or **spreading code**. The result resembles **white noise**. See Figure 7.

During de-spreading, the signal is again multiplied with the same spreading code. De-spreading thus converts the wideband signal into a narrowband one (this works due to the autocorrelation properties of the spreading code). At the same time, any narrowband interference is spread out.

Thus DSSS is more robust against (un)intentional interference and multipath effects, and narrowband jamming requires much more power. Broadband jamming is possible, but inherently requires much power.

Detecting DSSS signals is difficult, but not impossible (energy detection of strong signals, signal characteristics such as constant chip rate). Interception and modification is hard.

Example usages: GPS, 802.11b WiFi, CDMA (used in 3G). Non-military applications mainly use DSSS for interference-resistance and use public spreading codes. They are thus still vulnerable to malicious jamming as DoS.

Processing Gain PG Ratio of the spread bandwidth to the baseband bandwidth, in dB.




Figure 7: DSSS – baseband signal, spreading code, spread signal (top to bottom)

Chirp Signal / Sweep Signal Signal in which the frequency increases and decreases over time (“sweeping” over a bandwidth much wider than the baseband bandwidth). Narrowband and partial-band jamming are prevented, follower jamming not so much

Code-Division Multiple Access CDMA Multiple transmitters sending in the same area simultaneously, but using different spreading codes. Allows sharing of the same frequencies/bandwidth without interference.

3. Jamming-Resistant Broadcast

Broadcast Communication One sender, many receivers. Inherently open: receivers may join and leave at any time. All receivers listen (c.f. multicast). E.g. radio (FM/AM), GPS.

Challenges when securing broadcast:

many and unknown receivers, colluding receivers, internal + external attackers. In particular, plain spreading techniques (with group keys) do not work – an internal attacker can use their knowledge to jam other receivers.

Based on FHSS Broadcast Anti-Jamming System due to Desmedt et al.

Base station transmits on multiple frequencies simultaneously. Each receiver listens on a subset of frequencies at a given time. Protects against $j - 1$ colluding receivers, ensuring that each receiver has at least one non-jammed channel.

- **[Public] Channel Allocation Table:** Defines which channels any receiver should listen on, such that $j - 1$ receivers do not cover all channels of any other receiver (set coverage).
- **[Secret] Frequency Allocation Table:** Mapping from channel id to frequencies. Derived using a PRG. The complete table is only known to the base station.

Disadvantages: effectively a multicast solution since it requires a shared secret between the base station and each receiver.

Based on DSSS Dynamic Jamming Mitigation due to Chiang and Hu. Counteract jamming by using a balanced binary key tree.

Each node in the tree corresponds to a spreading code C_i . Each receiver N_i is assigned a leaf and knows all codes on the path from the root to that leaf.

The base station transmits on (a) a disjoint cover of codes (i.e. all users can decode exactly one code) and (b) a set of test codes. If a client receives a test code but not the detectable code, it reports jamming.

Splitting and reforming of the tree enables jamming mitigation.

Disadvantages: requires highly flexible, powerful base station. Requires a feedback channel. Requires a growing number of shared secrets (and receivers must be known).




Figure 8: Dynamic Jamming Mitigation – cover codes (red circle), test code (green square)




Figure 9: Circular dependency between anti-jamming and key establishment

Anti-Jamming—Key-Establishment Dependency Above techniques lead to a circular dependency. We need techniques without shared secrets! Idea: if we cannot coordinate sender and receiver, then don't even try (see next two techniques).

In addition, pre-loading shared keys in full of problems: requires a trusted party, key revocation, new clients joining, etc.

Uncoordinated Frequency Hopping Spread Spectrum UFH Neither attacker nor legitimate receivers can predict which channels are used. Equivalent to FH in terms of jamming protection (but not in throughput).

Transmitter steps:

1. Fragment message
2. Link fragments (against insertion)
3. Encode packets (ECC against jamming)
4. Repeated transmission while hopping on frequencies

Receiver steps: same process but reversed (plus packet ordering). Hops from one frequency to the other (sequentially is fine), in the hope of receiving a fragment.

Issue with fragment linking:

The signature is only verified at the end for the entire message.¹ Since there are exponentially many combinations for re-assembly, the attacker can now perform a DoS on a logical (rather than physical) level (pollution attack).

Solution: cryptographic linking of fragments (but without a shared key). E.g. hash linking, one-way accumulators, short signatures.

Disadvantages: Throughout up to 1000x less than FH. Higher latency (depending on attacker strengths, i.e. how high the chances are that the receiver gets a packet).

¹The signature is based on public-keys and a mutually trusted – but potentially offline – certificate authority CA.




Figure 10: UFH transmitter steps




Figure 11: UFH fragment linking – exponentially many candidate messages

Uncoordinated Direct Sequence Spread Spectrum UDSSS Neither attacker nor legitimate receivers can predict which spreading codes are used. The public code set C is composed of n code sequences, each containing l spreading codes.² Despreading is done by trial-and-error: it requires the correct code sequence and correct synchronisation (which fragment are we at?). The message is also repeatedly sent because of possible jamming – possibly in parallel to improve throughput.

Optimisation: first transmit the message M with a secret spreading code K using DSSS. Then transmit the spreading code K using UDSSS.

Advantage: quicker decoding, longer messages, flexible security level.

²This allows a message to be fragmented into l pieces.

4. Security of Global Navigation Satellite Systems GNSS

Overview Orbiting satellites transmit their location and a precise timestamp. Receivers collect these navigation messages and their arrival time and use **triangulation** to calculate their own position. Satellites are positioned such that at least four are always in sight on any point on Earth.

Three segments: users, satellites, ground control.³

Signalling Each satellite modulates the navigation message with a spreading code (coarse acquisition C/A for civilians (public), precision P/Y for military (secret)). The spread signal is then modulated onto a carrier.

Individual satellites use individual spreading codes to allow distinction.

GPS sends on two carrier frequencies at the same time, L1 ($1575.42 \text{ MHz} = 10.23 \text{ MHz} \times 154$) and L2 ($1227.60 \text{ MHz} = 10.23 \text{ MHz} \times 120$).⁴ Apart from jamming resistance and redundancy, this also allows to calculate the ionospheric delay error.

Due to atmospheric attenuation, down on Earth the GPS signal is well below the thermal noise.

Navigation message Each message consists of 25 frames. Each takes 30 sec to transmit, so the total time is 12.5 min.

Each frame contains: satellite clock + health data, 2x ephemeris (orbit details), other data + almanac (orbital + clock details).

Time of Arrival TOA Travel time of the signal from the satellite to the receiver. Used to calculate the distance and thus eventually the receiver position. Found by sliding the spreading code over the received message until a correlation peak.

Spoofing attacks Messages are unauthenticated (for practical reasons, else they would become too long).

By sending stronger signals, overshadowing the legitimate ones, an attacker can modify the *navigation message contents* (transmission time, satellite location) or their *time of arrival* (retransmitting captured signals with a temporal shift), resulting in a wrong location being calculated.

This is an issue in civilian GPS (messages can be generated and delayed) as well as in military GPS (messages can only be delayed since they are encrypted). Unfortunately, commercial GPS signal generators are becoming increasingly cheap.

4.1. Spoofing Detection and Mitigation

Types of countermeasures

- **Infrastructure/protocol:** e.g. cryptographic authentication of navigation messages
- **Receivers:** Use physical-layer characteristics of the signal to validate the signal as well as the calculated position/velocity/time. E.g. direction of arrival, carrier phase, signal strength, etc.

³There are of course issues with special and general relativity that mess with the time.

⁴This only applies to military. The civilian C/A is only transmitted on L1.




Figure 12: Seamless takeover attack

Angle of Arrival AoA Use multiple antennas (e.g. on both ends of a ship) to calculate the angle of arrival through the phase difference and the known distance between the antennas (see beam steering, Figure 4). In a spoofed scenario, the angles would all be very similar. Restricts the locations from which the attacker can successfully spoof.

Problems: Attacker can use drones to spoof signal from more realistic angle. Reflection of legitimate signal of buildings (thus reaching the receiver at a shallower angle) could be wrongly classified as spoofing. Computationally expensive phase measurement. Hardware modification.

Monitor Signal Characteristic Changes Over time, monitor signal properties such as AGC (Automatic Gain Control), noise level, number of satellites, spatial diversity (AoA) or the autocorrelation peak. Abrupt changes in any of these indicate presence of spoofing.

Seamless takeover attack The attacker starts transmitting a copy of a legitimate GPS signal in sync with the original one, but at low power, having no influence on the receiver. Then the attacker slowly starts increasing the power, until the receiver prefers the attacker signal. Now the attacker can change the GPS signal, and the receiver will keep following.

SPoofing REsistance GPS rEceiver SPREE Leverage peak tracking (of all signal peaks) to detect seamless takeover attacks. Navigation message inspection detects content spoofing.

Cryptographic approach (Kuhn)

1. At time t : satellite uses secret spreading code. Receiver uses a broadband receiver to capture the entire band.
2. At time $t+dt$: Satellite disclosed code, signing the disclosure with its secret key. Receiver verifies signature, de-spreads the signal.

Advantages: Prevents fake signal generation and individual signal delay.

Disadvantages: Requires pre-shared public satellite keys. Does NOT prevent full-band delay. Slightly inefficient (longer latency until signal lock). Replay attacks (?).

5. Secure Distance Measurement

Introduction

Applications: Wireless car keys, contact tracing in a pandemic, autonomous cars.

Attacks: **Replay attacks** are an issue, allowing an attacker to make devices appear physically closer (e.g. if the device naively use the observed signal strength to derive the distance).

Goals: (Provably) secure ranging, protecting against all logical and physical attacks and all attacker abilities. Focus on preventing distance reduction.

Current techniques (overview)

Non-Time-of-Flight:

- Received Signal Strength Indication RSSI (WiFi, Bluetooth, 802.15.4, NFC, RFID) – *insecure*
- (Multi-carrier) phase measurement⁵ – *insecure*
- Frequency-Modulated Continuous-Wave FMCW – *insecure*

Time-of-Flight:⁶

- Chirp Spread Spectrum (802.15.4 CSS) – *insecure*
- Ultra Wide Band UWB (802.15.4z) – *proposed*
- WiFi 802.11az – *efforts to secure OFDM-based*
- 5G – *first academic proposals*

Model On a logical level, we have a **verifier V** and a **prover P**, between which we want to measure the distance. A **malicious party M** attacks this.

See also the Brands-Chaum protocol (not discussed in HS20).

Additionally, we assume the worst case for the users but the best case for the attacker (bad channel/noise/multipath versus perfect channel → attacker guesses will seem like noise).




Figure 13: Model of the distance bounding scenario

⁵The distance is proportional to the phase.

⁶ Calculate via $d = c \cdot (t_{tof} - t_{proc})/2$ where t_{tof} is the time between sending and receiving the signal and t_{proc} is the known processing time on the responding device.

In general, manipulating time is harder than manipulating signal properties (strength, phase).

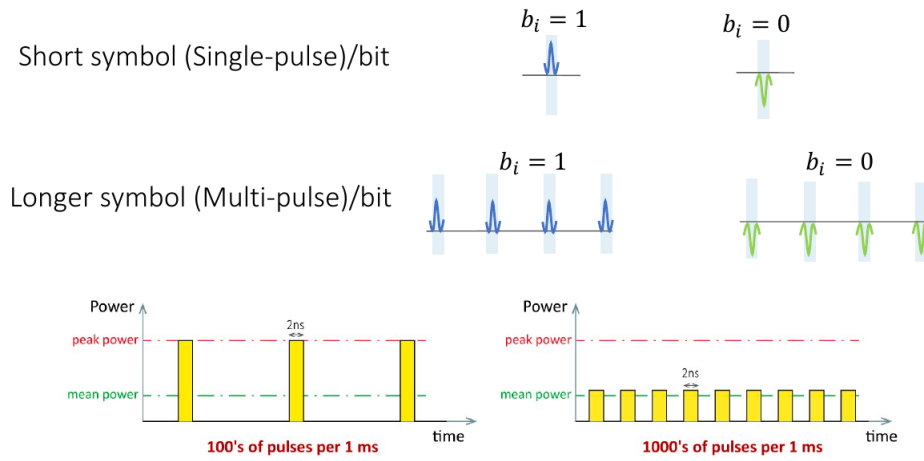


Figure 14: Bit representation: single- versus multi-pulse

Types of attacks (frauds)

- **Distance fraud:** A dishonest P tries to change its distance to V .
- **Mafia fraud:** Honest V, P being attacked by an external M .
- **Terrorist fraud:** Dishonest P and M collude to change P 's distance.
- **Distance hijacking:** Dishonest P leverages an honest P to change its distance.

Physical layer: Representing bits as pulses (UWB) There are two design options to represent a bit: either with a single strong pulse or a sequence of weaker pulses. Single pulses may not be detected reliably (distance, interference), so the aggregate over several pulses is the preferred representation.

Early-Detect/Late-Commit attack (ED/LC) This attacks shortens the distance, since the receiver receives the first symbols earlier than it should have.

1. Attacker sends noise (at time T_A)
2. Attacker learns correct symbol (at time T_{ed})
3. Attacker commits to correct symbol (at time T_{lc}), by sending the remaining pulses such that the sum over all pulses matches.

Note that this attack is not possible with single pulses. A single pulse is usually 1-2 ns long, so the attacker can cheat by at most 15-30 cm (performance/security tradeoff).

ED/LC Solution 1: Pulse Reordering UWB-PR Interleave pulses of subsequent symbols according to some cryptographic reordering. Thus the start and end time of a symbol is unpredictable, and the attacker can only guess.

The probability of an attack decreases [increases] with the number of interleaved bits [number of pulses per bit].

ED/LC Solution 2: Variance Based Detection Statistically analyse the received versus the expected pulses. This forces the attacker to “guess better” to reduce the variance and make their error indistinguishable from the noise.

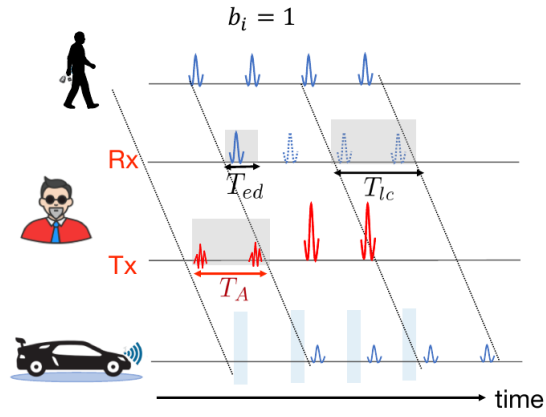


Figure 15: Early-Detect/Late-Commit attack (ED/LC)




Figure 16: Pulse Reordering

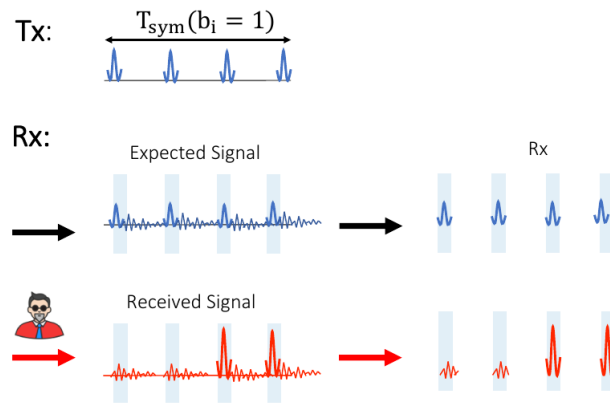


Figure 17: Variance Based Detection




Figure 18: Scrambled Timestamp Sequence

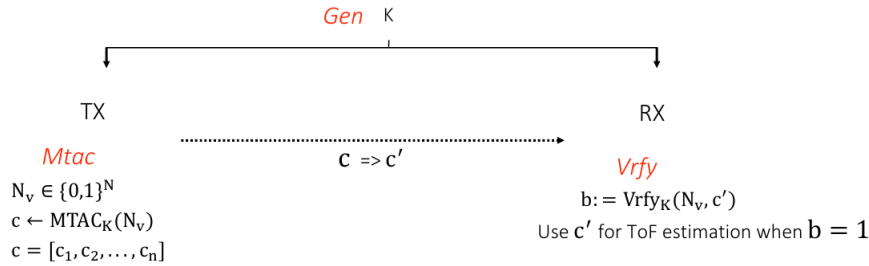


Figure 19: Message Time of Arrival Code MTAC

ED/LC Solution 3: Scrambled Timestamp Sequence After the preamble (high correlation) which is used for ToF *estimation*, send a Scrambled Timestamp Sequence STS (encrypted with a shared secret, low autocorrelation) to use for ToF *verification*. See IEEE 802.15.4z. Security not formally proven and unclear!

IEEE 802.15.4z LRP versus HRP

Low Rate Pulse LRP	High Rate Pulse HRP
Can use single pulse	No single pulse (energy too low)
Multi-pulse with UWB-PR efficient	UWB-PR + variance-based seem inefficient
Open security specs	No open security analysis
Low-cost, low-energy	

Message Time of Arrival Code MTAC New class of cryptographic primitives that verify the integrity of message arrival time. E.g.: single-pulse, UWB-PR, Variance-based detection.

Verifiable Multilateration Multiple verifiers with known locations want to determine the position of a prover. C.f. GPS trilateration. E.g. to position autonomous cars using cell towers.

Future Work

- Secure Positioning (e.g. verifiable multilateration)
- WiFi 802.11 and 5G ranging (some initial work)
- Efficient implementation + deployment

6. Security in Critical Transport Infrastructures

Air Traffic Control ATC

- **Primary Surveillance Radar PSR:** ground based, measures time delta between transmission and reflection \Rightarrow independent.
- **Secondary Surveillance Radar SSR:** Transponder based interrogation \Rightarrow dependent. Mode A (identification code), Mode C (identification code + barometric altitude), Mode S (selective addressing to interrogate a specific aircraft, used in *Traffic Alert and Collision Avoidance System*)
- **Automatic Dependent Surveillance-Broadcast ADS-B:** Aircraft determines its position via satellite and regularly broadcasts the result. Replaces functions of SSR, and enables inter-aircraft situational awareness.

Problem Statement Huge number of systems and protocols in aviation (see Figure 20). None has confidentiality, integrity or authentication. On the other hand, attacker capabilities grow as domain knowledge spreads and software defined radios become cheaper.

At the same time change is incredibly slow due to certifications and legacy compatibility.

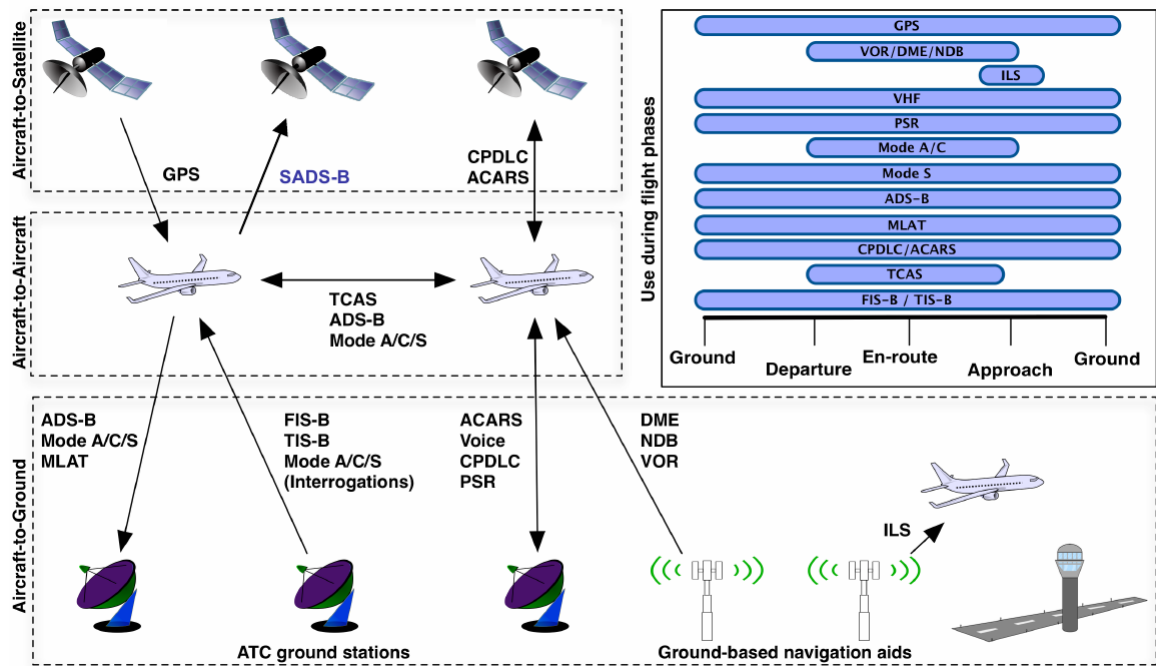


Figure 20: Overview of aviation systems

6.1. Privacy Issues in Aviation

No confidentiality of aircraft-ground communication Clear-text transmission of passenger medical status, forgotten belongings, credit card transaction details, etc.

Proprietary crypto ACARS⁷ datalink encryption uses a mono-alphabetic substitution cipher with a limited keyset hardcoded in a lot of private/military/government jets.

⁷ Aircraft Communications Addressing and Reporting System

Aircraft identifiers Aircraft transponders have unique IDs that include aircraft type and operator. These are not easily and legally changed (unless aircraft is sold).

No location privacy Aircraft are globally trackable by anyone. Either using websites like Flight-Radar24 (heavily filtered), ADS-B Exchange, OpenSky Network. Or collecting own data from SSR and ADS-B signals using a cheap radio.

Possible “uses” for tracking: government aircraft movement, mergers & acquisitions (M&A) activities.

Mitigations: Block aircrafts on tracking websites, obscure ownership (register to shell/trust companies), disable position broadcasts (still easily localised near departure/destination airports), use commercial transport.

6.2. Security Issues in Aviation

Attacks The usual candidates: jamming, modification, injection (ghost aircraft = DoS on ATC).

Safety vs Security *Safety* is about dealing with accidents and failures. We tackle it with experience (root cause analysis) and redundancy (decreasing the likelihood of failure of the entire system).

Security on the other hand is concerned with protecting against an intelligent, adaptive attacker. See the Swiss Cheese model in Figure 21.




Figure 21: Swiss Cheese Model: Safety vs Security

Traffic Collision Avoidance System TCAS Aircraft continuously predict each others' paths. If the paths are too close but not yet at risk, a *traffic advisory TA* is issued (and announced in the cockpit). If they remain on a close path, a *resolution advisory RA* is announced as compulsory instructions.

Attack: Attacker listens to target aircraft and injects TCAS responses, forcing a TA and RA, thus forcing the aircraft to make unwanted course changes (plus push pilots to reduce TCAS sensitivity).

6.3. Short-term Security Countermeasures

Cyber-Physical Countermeasures There won't be any real crypto any time soon. In the meantime, exploit physical layer data (timing, signal strength, Doppler shift, etc) to validate signals. Throw in some machine learning as well (for anomaly detection).

This hopefully lifts the bar back up to nation-state-attackers-only,

Multilateration Use multiple ground stations to receive (and validate) signals from aircraft.

OpenSky Network Crowdsources ATC. Volunteers install software-defined radios around the world to capture signals and publish them.

Short-term solution: rather than cryptographically authenticating messages between aircraft, use ground network to check if everybody else also received the same signal.

Advantages: Does not touch legacy systems, low cost, global coverage, flexible.

6.4. Satellite and Maritime Infrastructures

Satellite links Signals are receivable in a large area (continent-scale). This allows tracking of e.g. military aircraft from far away.

Maritime VSAT (Very Small Aperture Terminal). Still large and expensive. Connects ships to IP network on land (WiFi, fleet monitoring, weather, navigation, cargo, etc). Composed of a satellite uplink (large beam, since satellite needs to cover a large area \Rightarrow can be captured from far away) and a directed downlink (towards a ground/land station).

Example: *Electronic Chart Display and Information System ECDIS* (paper chart/map replacement) receives updates via VSAT.

Analysis: Lots of interesting yet unencrypted traffic – from standard DNS/VoIP/IMAP to specific “ship data”.

TLDR It's bad.

7. Physical Layer-based Security

7.1. Channel-based Key Establishment

Unique channels In a complex, multi-path rich environment channels exhibit time-varying, stochastic and reciprocal fading. For receivers separated by $> \lambda/2$ their channels are not correlated.

Thus the channels between sender S and receiver R are “random” and cannot be known/predicted by the attacker. In particular, the attacker cannot remotely measure multipath fading components of the signal strength.

Key Agreement through Channel Properties We can leverage different properties of the channel, e.g. RSSI⁸, CIR⁹ or signal phase.

The generic steps are the following:

1. Signal Acquisition and Quantisation
2. Reconciliation (error correction, privacy amplification)
3. Key confirmation

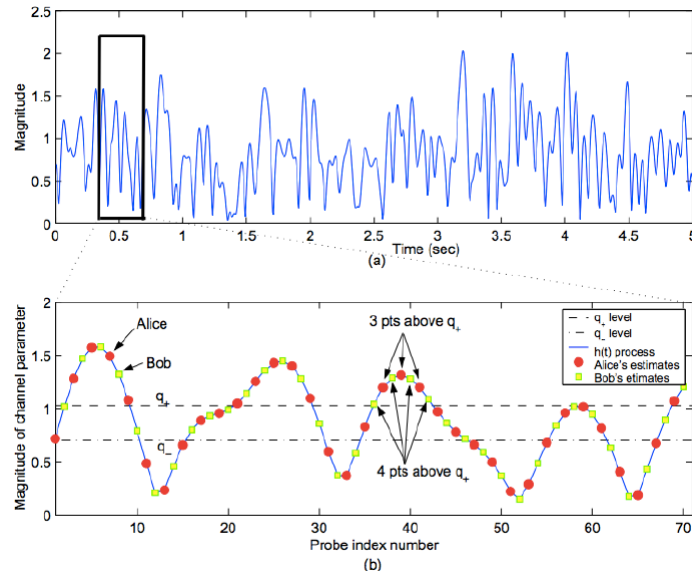


Figure 22: Measuring Channel Property Magnitudes over Time

Analysis There are several disadvantages to this scheme:

- No authentication: we establish a secret key, but with whom?
- No guarantees on the environment: Multi-path rich? Cannot be pre-measured? Receivers (and attacker!) at least $> \lambda/2$ from each other?
- Questionable benefit over classic public-private-key schemes (no information-theoretic security)
- Active attacker not considered: can influence and discover the key

⁸Received Signal Strength Indication

⁹Carrier-to-interference ratio




Figure 23: MIMO channel

Ensuring Secrecy with MIMO Goal: strengthen the security of the previous scheme.

Idea: use multiple antennas on both ends, to (a) steer the signal towards the receiver and away from the attacker and (b) use jamming to interfere with the attacker (but not the receiver).

Note that we can model each channel (or the signal on that channel) as a complex number (amplitude as the real part, phase as the imaginary part).

Zero Forcing Assumption: S knows the channels to the intended receiver R_1 and to the attackers R_2, R_3 , which are given as channel matrices H_r . The sender applies a transmission filter F that is constructed such that $R_r = HFD = HS$ contains useful data for R_1 but not for the attackers.

Orthogonal Blinding Same as zero forcing, but we only assume that the channel to the intended receiver R_1 is known (but not the channels to the attacker). Construct the filter F such that for everybody but R_1 the result R_r contains a jamming signal (i.e. noise).

Analysis

- ⊕ Stronger guarantees than SISO (beam forming focusses energy on receiver, jamming interferes with attacker)
- ⊖ Still no authentication
- ⊖ Still no guarantees on environment
- ⊖ Still questionable benefit over classic public-private-key schemes
- ⊖ Passive attack: known plaintext attack (attacker can train filter)
- ⊖ Active attack: abuse lack of authentication

7.2. Friendly Jamming for Confidentiality and Access Control

Friendly Jamming Transmit noise that the receiver can subtract (assuming a shared secret – possibly established with one of the earlier methods).¹⁰ The attacker on the other hand cannot distinguish signal and noise.

Example: IMD Shield for implanted medical devices.

¹⁰Note the difference to orthogonal blinding, which does not add noise to all channels but only to the null space of the receiver's channel.




Figure 24: Friendly Jamming Attack Through Antenna Placement

Analysis (as claimed by the IMD Shield paper).

Let DJ be the distance between the data and jamming antenna (of the sender). If $DJ > \lambda/2$ then an attacker with two antennas can separate the two signals (multiple channels). Furthermore, if $DJ \gg \lambda/2$ then the attacker can use directional antennas for signal separation.

Thus the only “safe” case is $DJ < \lambda/2$ (i.e. when channels from D to A and from J to A are highly correlated).

However, this does NOT hold: a MIMO attacker can retrieve the data even when $DJ < \lambda/2$, see the following:

Attack in Line of Sight LOS Model The attacker places two antennas A, B (see Figure 24). The ideal placement is such that they receive both the jamming signal simultaneously (e.g. equidistant to the jammer) and the data signals with a phase shift of $\lambda/2$. However, even with imperfect placement some data can still be recovered (with some attenuation).

When we also consider multipath effects (additional changes in amplitude and phase offsets) the attack becomes much harder.

Conclusion Jamming works well for access control (in the sense that J prevents malicious signals being sent to D). It does NOT work for confidentiality since a MIMO attacker can retrieve data even when $DJ < \lambda/2$.

In other words, the IMD Shield protects against malicious commands being sent to the implant, but not against an attacker listening to outgoing signals.

Signal Annihilation An attacker can deliberately introduce destructive interference to attenuate a legitimate signal.

Note that this is different to jamming! Here, we are fooling the receiver to believe that there is no signal on the channel.

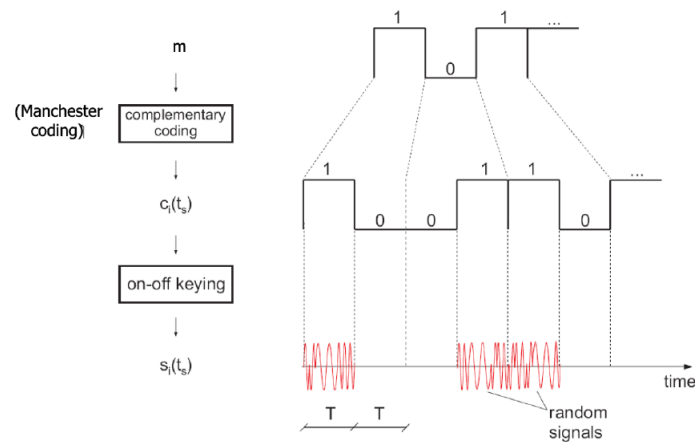


Figure 25: Integrity Codes

Summary

- Using the physical layer for access control seems realistic.
- Using it for confidentiality however is questionable! Weak guarantees, use only as complementary measures.

7.3. Broadcast Authentication: Presence Awareness

Scenario One broadcasting station, many receivers. No pre-shared keys, no credentials (certificates, public keys, etc). Receivers know that they are within the power range of the legitimate sender (reasonable in airports, universities, etc). Receivers know which channel (frequency) the sender is broadcasting. The sender is always on and transmitting.

Goal: distinguish between broadcast messages of legitimate and malicious station. In other words, a physical layer-based broadcast authentication scheme based on presence awareness. See also subsection 8.1.

Integrity Codes Protocol

Sender: Spread message m from k to $2k$ bits using Manchester encoding ($1 \rightarrow 10$, $0 \rightarrow 01$). Transmit result using on-off keying.

Receiver: Set power thresholds above which to interpret a signal as 1. Then decode.

Integrity Verification: Check if hamming weight¹¹ $H(m) = |m|/2$. If yes, the message was not modified.

Analysis

- The attacker can easily change $0 \rightarrow 1$ in the raw/keyed signal. They can change $1 \rightarrow 0$ only with small probability (assuming signal annihilation is hard).¹²
- How to handle arbitrary length messages? Solution: Use i-delimiters between messages. In Manchester encoding e.g. 111000.
- Slow. Solution: broadcast hash of message using integrity codes, use another faster channel for the full message.

¹¹Recall that the hamming weight is the number of bits unequal to 0.

¹²thgoebel: Why is this so? Just a couple of slides ahead it was explained that signal annihilation is doable?

8. Broadcast Authentication and Device Pairing

8.1. Broadcast Authentication: Delayed Key Disclosure

Scenario One sender, many unknown (possibly malicious) receivers. All (legitimate) receivers need to verify the authenticity of the broadcast messages.

Challenge: don't use public key cryptography (computationally expensive, especially on low-power devices), but instead only rely on symmetric cryptography.

One-way hash chain Repeated application of a hash function F , starting at an original value s_l (see Figure 26). Due to the one-way property of hash functions an attacker knowing s_i can only compute the subsequent value in the chain (s_{i-1}) but not the value s_{i+1} that generated s_i . This allows us to “use” the values s_i one-by-one and always reveal a single new value. See also Merkle hash trees.




Figure 26: One-way Hash Chain

Delayed Key Disclosure (TESLA) See Figure 27.

- The sender randomly generates a secret K_l , computes the entire hash chain and distributes K_0 as a “public key” to all receivers.
- To send a message M_j in the interval i , it sends

$$P_j = \{M_j || MAC(K'_i, M_j) || K_{i-d}\}$$

where $||$ denotes concatenation and usually $d = 1$. Note that we can send multiple messages in the same time interval using the same key.

- To verify a message, the receiver needs to receive M_j and – at a later time step and in another message – the key K_i . They can then compute K'_j .¹³ Then the receivers verify whether the MAC matches and the key K_i does correctly hash down to the pre-loaded K_0 and whether the message was received in an interval where the key was valid.
- If a key is used after the interval, the message is ignored.

Analysis:

TESLA achieves asymmetry by delaying the explicit disclosure of the self-authenticating keys in clear-text. It requires time to be loosely synchronised between the sender and receivers.

8.2. Device Pairing

Scenario We want to establish a secret key between two wireless devices, in the presence of an adversary.

¹³Why this extra hashing step is needed is “an implementation detail”.

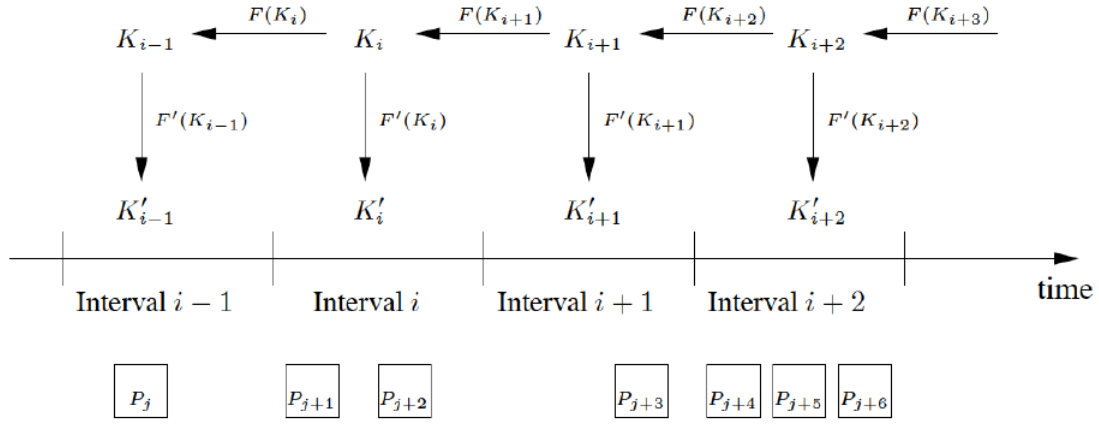


Figure 27: TESLA




Figure 28: Diffie-Hellman Key Exchange and MITM Attack

Diffie-Hellman Key Exchange The classic secret key establishment scheme, established a shared key $k_{AB} = g^{ab} \bmod p$. See Figure 28 for a refresher.¹⁴ However, it is only secure against a passive adversary: an active attacker can trivially MITM the key exchange.

Goal: Authenticate the DH key exchange between two wireless devices, to ensure they have established the same key.

Device Pairing Proposals A selection:

- Short string comparison: hash the key, display the string on both devices.
- Seeing is Believing: One device scans a QR code of the other's public key (one-way authentication).
- Load and Clear: has public key, map it to a recognisable sentence, read it out loud (text-to-speech TTS).
- Integrity Regions: use distance bounding to authenticate the key if the devices are in close proximity.
- Resurrecting duckling: first to receive the key becomes the owner for life.
- Shake them up: see below.
- PIN/Passkey entry: Bluetooth.

¹⁴We have prime p , a generator g of Z_p^* and the Diffie-Hellman assumption stating that the discrete logarithm is hard in some groups.

Note the different security assumptions: only friends can be close, trust on first use, trusting the human, binding the success of the authentication to the happening and success of authentication (eliminating human error).

Shake them up

Idea:

Divide time into N slots. In each slot, randomly either A or B transmit a message. A message is either 1 ("I am Alice") or 0 ("I am Bob"). Depending on who sent the message, this statement is either true or false. Of course, both Alice and Bob know which is the case. If the statement is true, set the next bit-to-be-exchanged to 1 otherwise to 0, thus creating a shared secret.

Synchronisation and key exchange are initiated through physically shaking the device (hence the name).

Analysis:

Assumes that Eve is too far away to distinguish messages from A and B , and can thus not know whether a statement was true or false. This assumption may NOT hold as it can be attacked using signal fingerprinting to distinguish the source of a signal.

Eve can insert messages, but then A and B will set their bits to different values (since none sent the message and will assume a false statement), thus only creating a DoS situation.




Figure 29: Shake Them Up

9. WiFi Security

9.1. WiFi Basics

Terminology

- *Station (STA)/client*: terminal with access to the wireless media
- *Access Point (AP)*: Station integrated both with the wireless media and the distribution system
- *Service set/extended service set ESS*: Group of nodes logically grouped together and identified by their SSID
- *Basic service set (BSS)*: Subgroups of a service set using the same radio frequency/the same AP (or in general: the same physical layer medium access means).
- *Service set identifier (SSID)*: 32 byte name (usually human-readable) to identify a network
- *Channel*: 20 MHz wide frequency range to use for WiFi, typically around 2.4 GHz and 5 GHz (see Figure 2).
- *Medium Access Control (MAC)*: Goal: deliver data reliably and securely while sharing the open medium




Figure 30: WiFi System

Carrier-Sense Multiple Access with Collision Avoidance CSMA/CA Medium access control mechanism. Part of the *Distributed Coordination Function DCF*.

1. *Carrier sense* – monitor medium to check if it is idle
2. *Collision avoidance* – if another node sends, wait for a randomised *backoff period*, then listen again
3. *Transmit* – send entire frame, wait for ACK, if no ACK then backoff and wait

Standard	802.11a	802.11b	802.11g	802.11n	802.11ac	802.11ad
Year introduced	1999	1999	2003	2000	2012	2014
Maximum data transfer speed	54 Mbps	11 Mbps	54 Mbps	65 to 600 Mbps	78 Mbps to 3.2 Gbps	6.76 Gbps
Frequency band	5 GHz	2.4 GHz	2.4 GHz	2.4 or 5 GHz	5 GHz	60 GHz
Channel bandwidth	20 MHz	20 MHz	20 MHz	20, 40 MHz	40, 80, 160 MHz	2160 MHz
Highest order modulation	64 QAM	11 CCK	64 QAM	64 QAM	256 QAM	64 QAM
Spectrum usage	OFDM	DSSS	DSSS, OFDM	OFDM	SC-OFDM	SC, OFDM
Antenna configuration	1×1 SISO	1×1 SISO	1×1 SISO	Up to 4×4 MIMO	Up to 8×8 MIMO, MU-MIMO	1×1 SISO

Figure 31: WiFi Standard Versions




Figure 32: Hidden Terminal Problem

Hidden terminal problem Occurs if a node B can communicate with two nodes A, C that can not communicate with each other. That is, if they both tried to send data to B they would individually sense the medium to be idle, but then they would clash at B anyway. See Figure 32.

The solution is to send a *Request-to-Send RTS* message before transmitting a frame and wait for all nodes to reply with a *Clear-to-Send CTS*. Now when A sends a RTS even though C does not receive it, C still receives the CTS from B and will thus back off.

This feature is optional in IEEE 802.11!

Frame types Data frames (user traffic), control frames (ACK, RTS/CTS), management frames (beacon¹⁵, association request, deauthentication).

¹⁵Contains basic information about the network, e.g. SSID, FH parameters, DSSS parameters, other supported options.

9.2. Basic Manipulations

Capturing all frames Easy – just install any packet capture software (such as Wireshark) and set the wireless interface to “promiscuous mode”.

Communication (un)fairness DCF is fair – assuming all stations adhere to it! However by modifying backoff parameters in the wireless driver one can behave selfishly to get higher throughput.

Jamming

Simple:

- (a) Trigger carrier sense of transmitter (prevent sending, less power).
- (b) Mangle the frame at the receiver (prevent receiving, more power).

Selective:

Listen, decode prefix of an incoming frame and based on recipient address decide whether to jam the remaining data portion of the frame. One needs to be very fast!

Man-in-the-middle Clone real AP onto different channel → get clients to sign up to yours (e.g. jam real AP, send a Channel Switch Announcement CSA) → inspect/manipulate/forward frames to real AP → profit.

9.3. WiFi Security Standards

Wired Equivalent Privacy WEP

First WiFi security standard (1997), with the goal of making wireless as secure as wired networks. Today fully broken.

Operation:

Checksum: plaintext $P = (m, CRC(m))$

Encryption: ciphertext $C = P \oplus RC4(IV, k)$ (stream cipher RC4 with shared key k , initialisation vector IV)

Transmit (IV, C) .

Problems:

- *Confidentiality*: Keystream reuse allows recovery of plaintexts:

$$C_1 \oplus C_2 = (P_1 \oplus RC4(IV, k)) \oplus (P_2 \oplus RC4(IV, k)) = P_1 \oplus P_2$$

Since the IV has 24 bits, at a data rate of 5 Mbps it will repeat after \approx half a day. Knowing one of the plaintexts is reasonable (common structures such as IP headers, packet injection from the internet, other redundancies).

- *Integrity*: CRC-32 is not a cryptographic MAC. This allows controlled message modification, i.e. given a $(ciphertext, checksum)$ pair we can create another $(ciphertext', checksum')$ pair such that the checksum is valid for the decrypted $plaintext'$.
- *Access control*: works as follows:
 1. AP sends a plaintext challenge
 2. Station replies with WEP encryption of challenge (proving possession of the key)

3. AP completes association request

Through simple traffic capture an adversary can learn a valid plaintext/ciphertext pair. From that they can derive the keystream to later compute a valid response to a new challenge (since the station can choose the IV).

Since 2007, tools like aircrack-ng allow recovering WEP keys in a matter of minutes.

Wrong idea during design: “in-flight” modification of packets may be hard, but don’t forget the possibility of offline attacks!

WiFi Protected Access WPA / Temporal Key Integrity Protocol TKIP

Designed as a transitional mechanism in 2003 and to be compatible to WEP devices (though in 2019 still $\approx 50\%$ of networks accept it...).

Idea:

Prevent keystream reuse and use a cryptographic MAC, while staying compatible to WEP. “Achieved” by augmenting encryption with per-packet key mixing, RC4 keystream filtering, a new integrity mechanism (MICHAEL) and counters for replay protection (TSC). This allows to keep the WEP mechanism in hardware but change the inputs given to it.

ChopChop attack: (for WEP!)

Capture an encrypted WEP frame. Remove the last data byte, guess it, re-compute checksum for the shorter packet, use AP as an oracle (must discard frames whose checksums fail). Rinse and repeat to learn one data block after the other.

For a detailed explanation see this aircrack doc entry.

Since this is a stream cipher, knowing one plaintext/ciphertext pair also reveals a keystream.

A similar attack also works for WPA/TKIP since MIC(MICHAEL) is reversible and the AP reports failures.

Biased keystream:

There exist statistical biases in a RC4 keystream. With enough packets this can be used to recover the plaintext for one MIC key. Computationally expensive, but shows that crypto can be attacked, too.

Summary:

Difficult start (had to reuse WEP hardware). Raised the bar (despite ugly fix), but attacked in 2009, 2014/2015. Okayish transition mechanism.

WPA2

Introduced in 2004. Better encryption (AES-128) and integrity (CBC-MAC, using *authenticate-then-encrypt*). Better authentication with a 4-way handshake.

Handshake:

Goal: Starting from a pre-shared key (PMK) provide mutual authentication and establish a session key (PTK).

Steps: exchange nonces (msg1,2), derive PTK using PMK + nonces + MAC addresses, authenticate exchange (msg3,4).

Proven “secure” in 2005 (mutual authentication and password is not leaked).

Key Reinstallation Attack (KRACK) (2017):

Goal: force a keystream reuse.

Observations: AP may retransmit msg3 if no ACK. Each time, client reinstalls the *same* PTK, and

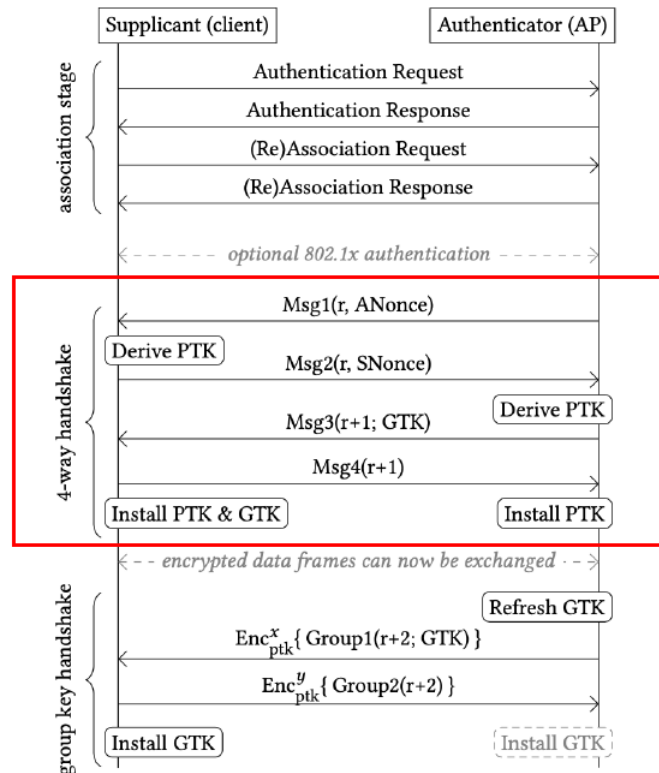


Figure 33: WPA2 Handshake

in the process it resets some counters. This resulting in the same keystream being reused.

Approach: Replay msg3.

Summary:

Solid cryptography/encryption, weakness in handshake. Proven properties hold, but model does not capture *when* a key is installed. Can be patched (and re-attacked).

WPA3

Introduced in 2018. Updated cryptography (AES-128/256 encryption, SHA-384 HMAC integrity), new handshake (based on the Dragonfly Key Exchange (RFC 7664)).

Handshake improvements:

Turns a low-entropy password into a high-entropy key, thus allowing for shorter passwords. Has forward secrecy.

Dragonblood (2020):

Transition mode allows downgrade to WPA2.

Timing attack: execution time of turning the password into a group element at the start of the handshake depends on the password (i.e. a side-channel attack, though only for bad curves).

A. Imprint

This document closely follows the lecture slides of the *Security of Wireless Networks* lecture in the autumn semester 2020 at ETH Zurich. Our contribution to this is editing the whole lot and refactoring even more so that it may fit the "lecture summary" style. However, basically all graphics are copy & pasted from the slides. If you don't want yours here, please contact us and we will remove them.

In addition, this summary is based on a summary by Sarah Kamp.

Otherwise, our part of the work is published as CC BY-NC-SA.